

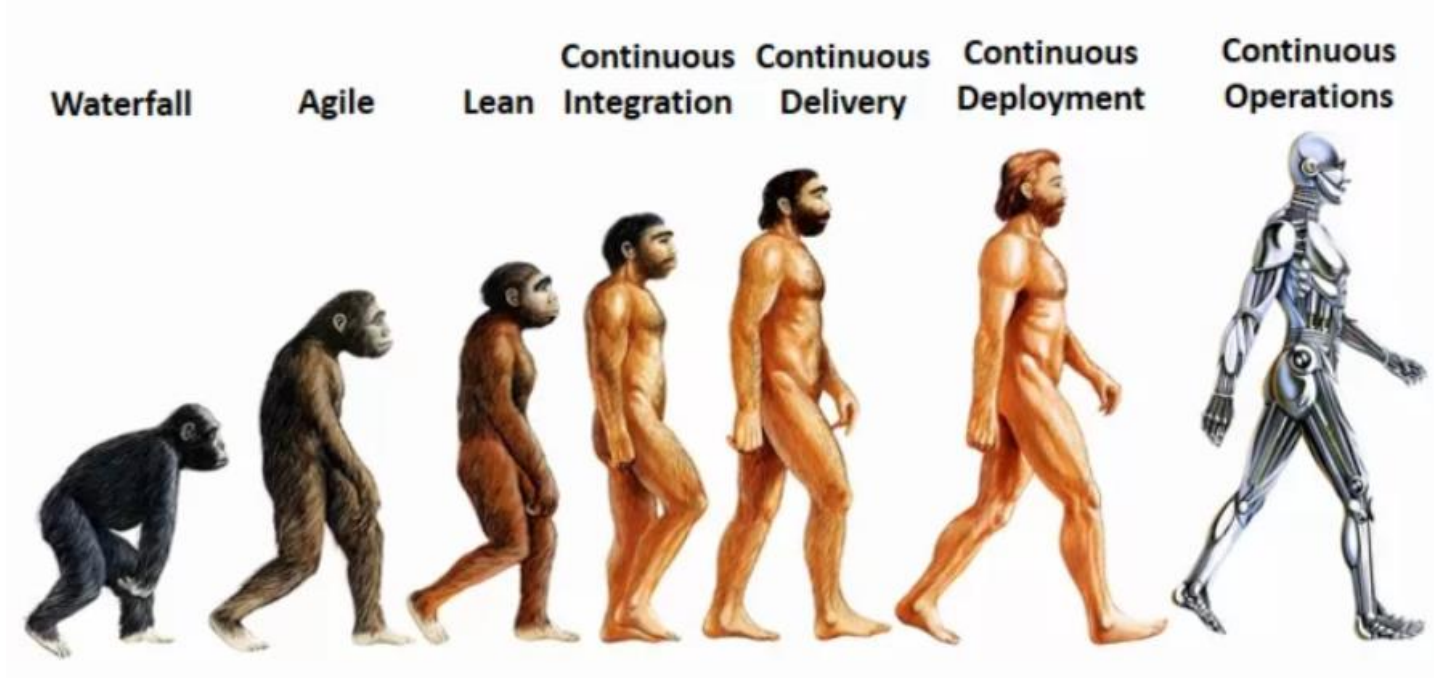
docker

Docker Fundamentals

The open platform to build, ship and run any application anywhere

About me

- 14+ Engineering experience (Cisco Systems, VMware)
- Continuous Operations freak and AWS enthusiast



Session Logistics

- 3 hours (exercising time included)
- “Introduction to Docker” course required
- You will need:
 - Linux machine
 - Dedicated
 - VM
 - Cloud provided
 - Enthusiasm
- Registration at hub.docker.com

Recap from “Introductions to Docker”


- Intro to Docker
- Benefits of Container based virtualization
- Docker concepts and terms
- Simple container examples
- Commands
 - `docker run`
 - `docker ps`
 - `docker images`

Agenda

- Building Images
- Dockerfile
- Managing Images and Containers
- Pushing Images on Docker Hub
- Docker Volumes
- Basic Container networking
- Configuring our first app


Image Layers

- Images are comprised of multiple layers
- A layer is also just another image
- Every image contains a base layer
- Docker uses a copy on write system
- Layers are read only



The Container Writable Layers

- Docker creates a top writable layer for containers
- Parent images are read only
- All changes are made at the writeable layer




Docker Commit

- **docker commit** command saves changes in a container as a new image
- Syntax
`docker commit [options] [container ID] [repository:tag]`
- Repository name should be based on username/application
- Can reference the container with container name instead of ID

Build New Image

1. Create a container from an Ubuntu image and run a bash terminal
`docker run -i -t ubuntu:14.04 /bin/bash`
2. Inside the container, install curl
`apt-get install curl`
3. Exit the container terminal
4. Run `docker ps -a` and take note of your container ID
5. Save the container as a new image. For the repository name use <your name>/curl. Tag the image as 1.0
`docker commit <container ID> <yourname>/curl:1.0`
6. Run `docker images` and verify that you can see your new image

Dockerfile



Intro to Dockerfile

*A **Dockerfile** is a configuration file that contains instructions for building a Docker image*

- Provides a more effective way to build images compared to using `docker commit`
- Easily fits into your continuous integration and deployment process

Dockerfile Instructions

- Instructions specify what to do when build the image
- FROM instruction specified what the base image should be
- RUN instruction specified a command to execute

```
#Example of a comment  
FROM ubuntu:14.04  
RUN apt-get install vim  
RUN apt-get install curl
```

Run Instructions

- Each RUN instruction will execute the command on the top writable layer and perform a commit of the image
- Can aggregate multiple RUN instructions by using “&&”

```
RUN apt-get update && apt-get install -y \  
    curl \  
    vim \  
    openjdk-7-jdk
```

Docker Build

- Syntax
`docker build [options] [path]`  **build context**
- Common option to tag the build
`docker build -t [repository:tag] [path]`

Build an image using the current folder as the context path. Put the image in the johnnytu/myimage repository and tag it as 1.0

```
docker build -t johnnytu/myimage:1.0 .
```

As above but use the myproject folder as the context path

```
docker build -t johnnytu/myimage:1.0 myproject
```


Build from Dockerfile

- Build your own image and verify the **curl** is inside the container

EXERCISE

CMD Instructions

- CMD defines a default command to execute when a container is created
- CMD performs no action during the image build
- Shell format and EXEC format
- Can only be specified once in a Dockerfile
- **Can be overridden at run time**

Shell format

```
CMD ping 127.0.0.1 -c 30
```

Exec format

```
CMD ["ping", "127.0.0.1", "-c", "30"]
```


Try CMD

1. Go into the test folder and open your Dockerfile from the previous exercise
2. Add the following line to the end
`CMD ["ping", "127.0.0.1", "-c", "30"]`
3. Build the image
`docker build -t <yourname>/testimage:1.1 .`
4. Execute a container from the image and observe the output
`docker run <yourname>/testimage:1.1`
5. Execute another container from the image and specify the echo command
`docker run <yourname>/testimage:1.1 echo "hello world"`
6. Observe how the container argument overrides the CMD instruction

ENTRYPOINT Instruction

- Defines the command that will run when a container is executed
- Run time arguments and CMD instruction are passed as parameters to the ENTRYPOINT instruction
- Shell and EXEC form
- EXEC form preferred as shell form cannot accept arguments at run time
- Container essentially runs as an executable

```
ENTRYPOINT ["ping"]
```

Start and Stop Containers

List all containers

```
docker ps -a
```

Start a container using the container ID

```
docker start <container ID>
```

Stop a container using the container ID

```
docker stop <container ID>
```

Getting terminal access

- Use **docker exec** command to start another process within a container
- Execute `/bin/bash` to get a bash shell
- `docker exec -i -t [container ID] /bin/bash`
- Exiting from the terminal will not terminate the container

```
johnnytu@new-docker:~$ docker exec -it serene_shockley bash
root@4cfbac7ba80a:/usr/local/tomcat# cd
root@4cfbac7ba80a:~# ps -ef
UID          PID    PPID    C  STIME TTY          TIME CMD
root           1         0   3  04:57 ?           00:00:03 /usr/bin/java
root          34         0   0  04:59 ?           00:00:00 bash
root          40        34   0  04:59 ?           00:00:00 ps -ef
root@4cfbac7ba80a:~#
```

Clean-up

- Containers

- `docker rm`

- `docker rm $(docker ps -a -q)`

- Images

- `docker rmi`

- `docker rmi $(docker images -q --filter "dangling=true")`

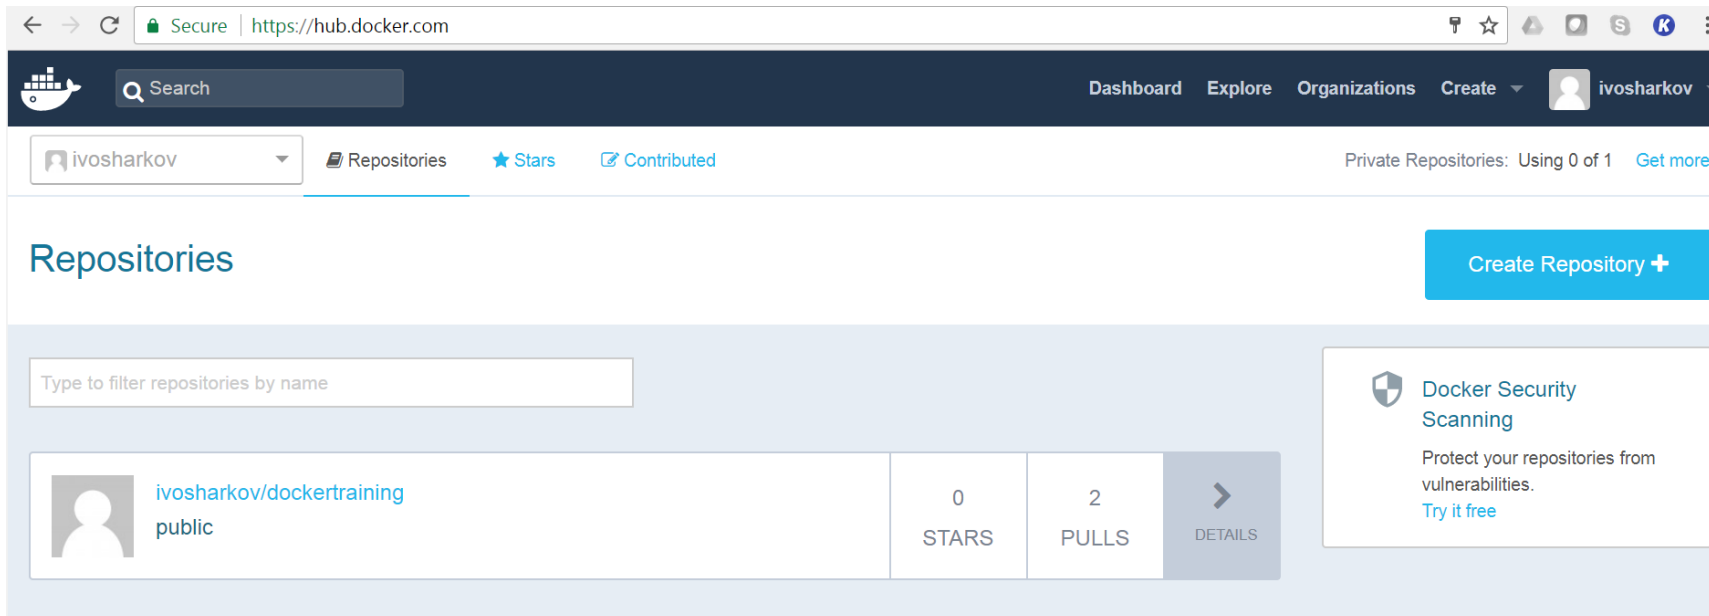
- Volumes

- `docker volumes rm`

- `docker volume rm $(docker volume ls -qf dangling=true)`

Docker Hub Repositories

- Users can create their own repositories on Docker Hub
- Public vs Private
- Push local images to a repository



Pushing Images to Docker Hub

- Use **docker push** command
- Syntax: **docker push repo:tag**
- Local repo **MUST** have the same name and tag as the Docker Hub repo

Tagging Images

- Used to rename a local image repository before pushing to Docker Hub
- Syntax:
`docker tag [image ID] [repo:tag]`
OR
`docker tag [local repo:tag] [Docker Hub repo:tag]`

Tag image with ID (trainingteam/testexample is the name of repository on Docker hub)

```
docker tag edfc212de17b trainingteam/testexample:1.0
```

Tag image using the local repository tag

```
docker tag johnnytu/testimage:1.5 trainingteam/testexample
```


Push to Docker Hub

1. Login to your Docker Hub account
2. Create a new public repository called “testexample”
3. Tag your local image to give it the same repo name as the repository you created on Docker Hub

```
docker tag <yourname>/testimage:1.1  
<yourname>/testexample:1.1
```

4. Push the new image to Docker Hub

```
docker push <yourname>/testexample:1.1
```
5. Go to your Docker Hub repository and check for the tag

Volumes

*A **Volume** is a designated directory in a container, which is designed to persist data, independent of the container's life cycle*

- Volume changes are excluded when updating an image
- Persist when a container is deleted
- Can be mapped to a host folder
- Can be shared between containers

Mount a Volume

- Volumes are mounted when creating or executing a container
- Can be mapped to a host directory
- Volume paths specified must be absolute

Execute a new container and mount the folder /myvolume into its file system

```
docker run -d -P -v /myvolume nginx:1.7
```

Execute a new container and map the /data/src folder from the host into the /test/src folder in the container

```
docker run -i -t -v /data/src:/test/src nginx:1.7
```

Volumes in Dockerfile

- VOLUME instruction creates a mount point
- Can specify arguments JSON array or string
- Cannot map volumes to host directories
- Volumes are initialized when the container is executed

String example

```
VOLUME /myvol
```

String example with multiple volumes

```
VOLUME /www/website1.com /www/website2.com
```

JSON example

```
VOLUME ["myvol", "myvol2"]
```

Uses of volumes

- De-couple the data that is stored from the container which created the data
- Good for sharing data between containers
 - Can setup a data containers which has a volume you mount in other containers
- Mounting folders from the host is good for testing purposes but generally not recommended for production use

Create and test a Volume

1. Execute a new container and initialise a volume at /www/website.
Run a bash terminal as your container process
`docker run -i -t -v /www/website ubuntu:14.04 bash`
2. Inside the container, verify that you can get to /www/website
3. Create a file inside the /www/website folder
4. Exit the container
5. Commit the updated container as a new image called test and tag it as 1.0
`docker commit <container ID> test:1.0`
6. Execute a new container with your test image and go into it's bash shell
`docker run -i -t test:1.0 bash`
7. Verify that the /www/website folder exists and that there are no files inside

Mapping ports

- **Recall:** containers have their own network and IP address
- Map exposed container ports to ports on the host machine
- Ports can be manually mapped or auto mapped
- Uses the `-p` and `-P` parameters in **docker run**

Maps port 80 on the container to 8080 on the host

```
docker run -d -p 8080:80 nginx:1.7
```

Automapping ports

- Use the `-P` option in `docker run`
- Automatically maps exposed ports in the container to a port number in the host
- Host port numbers used go from 49153 to 65535
- Only works for ports defined in the EXPOSE instruction


Auto map ports exposed by the NGINX container to a port value on the host

```
docker run -d -P nginx:1.7
```


Linking Containers

***Linking** is a communication method between containers which allows them to securely transfer data from one to another*

- Source and recipient containers
- Recipient containers have access to data on source containers
- Links are established based on container names



Create a Link

- Create the source container first
- Create the recipient container and use the **--link** option
- Best practice – give your container meaningful names

Create the source container using the postgres

```
docker run -d --name database postgres
```

Create the recipient container and link it

```
docker run -d -P --name website --link database:db nginx
```

Linking Use cases

- Containers can talk to each other without having to expose ports to the host
- Essential for micro service application architecture
- Example:
 - Container with Tomcat running
 - Container with MySQL running
 - Application on Tomcat needs to connect to MySQL

Link two Containers

1. Run a container in detached mode using the `postgres` image.
Name the container “dbms”

```
docker run -d --name dbms postgres
```
2. Run another container using the `Ubuntu` image and link it with the “dbms” container. Use the alias “db”, run the `bash` terminal as the main process

```
docker run -it --name website --link dbms:db ubuntu:14.04 bash
```
3. In the “website” container terminal, open the `/etc/hosts` file
4. What can you observe?

Make hands dirty

- Start Nginx `-p 555:80` and Jenkins `-p 8000:8080` images with persistent volumes on the file system
- Ensure you are able to login in `localhost:8000` and the data will endure `docker rm -f jenkins`
- Modify Nginx configuration so any request to `http://localhost:555/jenkins` should be redirected to our Jenkins service

Questions



References

- Dockerfile best practices

https://docs.docker.com/engine/userguide/eng-image/dockerfile_best-practices/

- CMD vs Entry

<https://www.ctl.io/developers/blog/post/dockerfile-entrypoint-vs-cmd/>

- Books for free

<http://gen.lib.rus.ec/>

- Google 😊